

Week 6 – Lesson 2: Long Short Term Memory (LSTM) Networks

Dr. Hongping Cai
Department of Computer Science
University of Bath




Topic 1: LSTM-1

Standard/Vanilla RNNs


- In a standard RNN, repeating modules contain a **simple computation** node.

$$h_t = \tanh\left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_h\right)$$




LSTM

- LSTMs are explicitly designed to deal with the long-term dependency problem.
- LSTM modules contain computational blocks that control information flow.



Two core concepts of LSTMs

- **Gates:** to control what information is to keep and forget.
- **Cell state:** act as a transport highway that transfers relative information all way down the sequence chain, thus store long-term information.



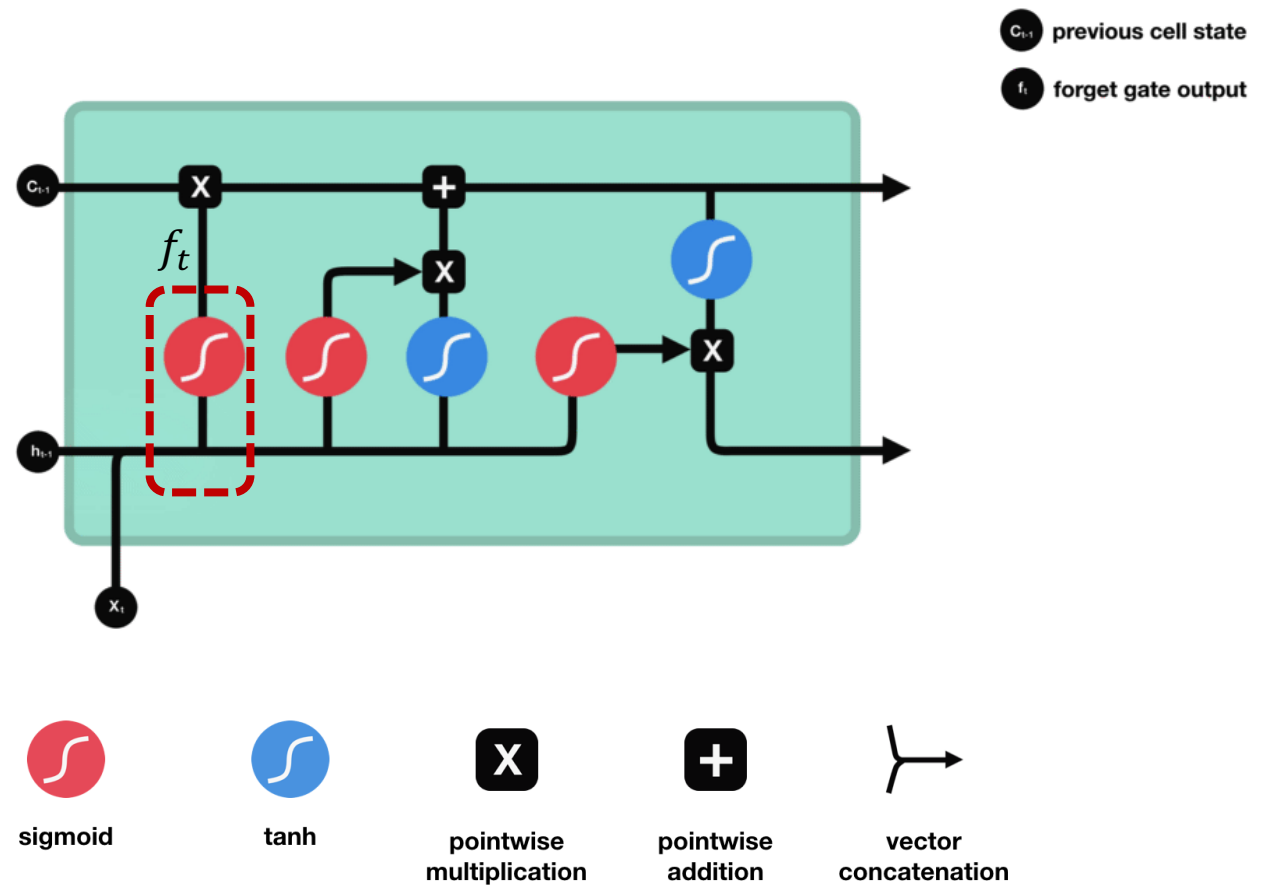
All gate values are between 0 (**discard**) and 1 (**keep**).



Forget gate

- **Forget gate:** forget irrelevant parts of the previous state.

$$f_t = \sigma(W_f \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_f)$$



Animation from: <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

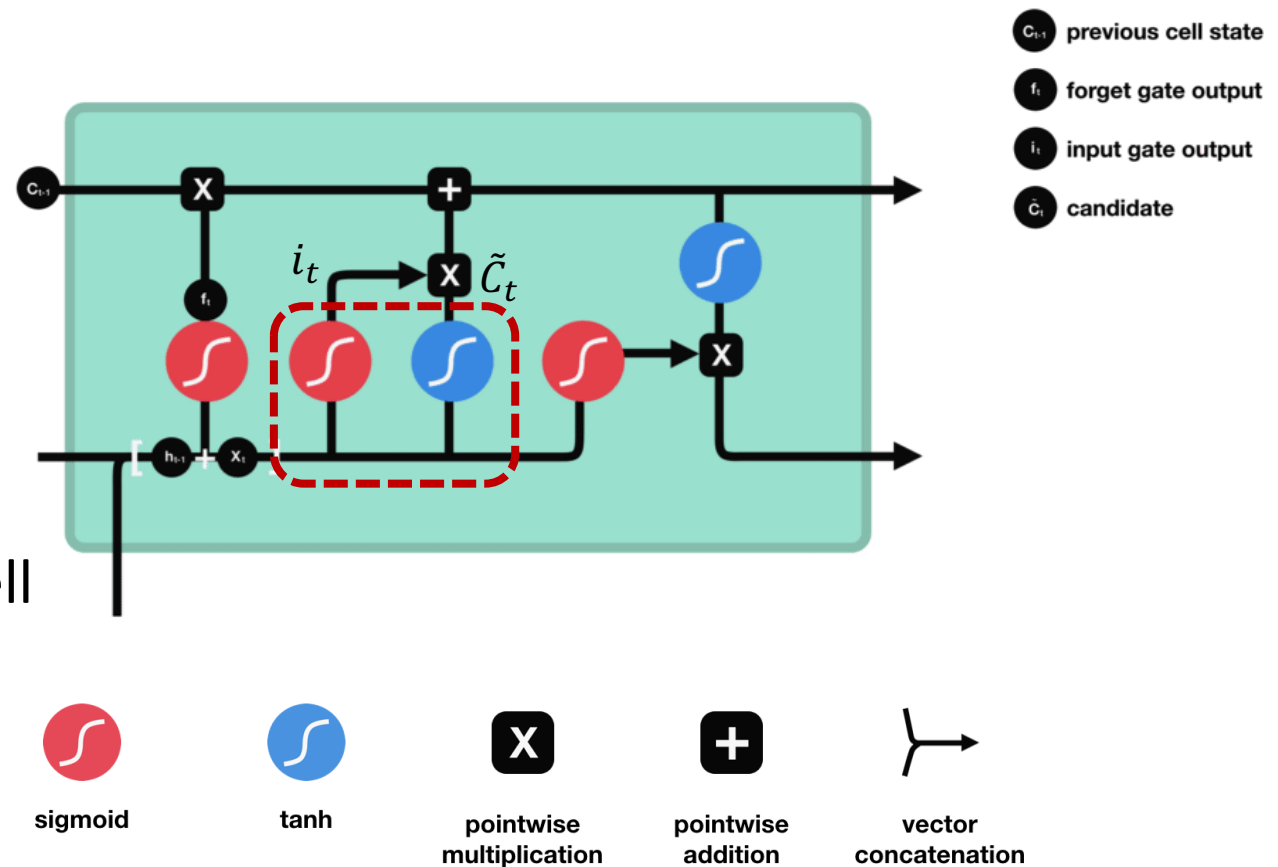
Input gate

- **Input gate:** decides what information is relevant to add from the current step.

$$i_t = \sigma(W_i \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_i)$$

- **New cell content:** the new content to be written to the cell

$$\tilde{c}_t = \tanh(W_c \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_c)$$



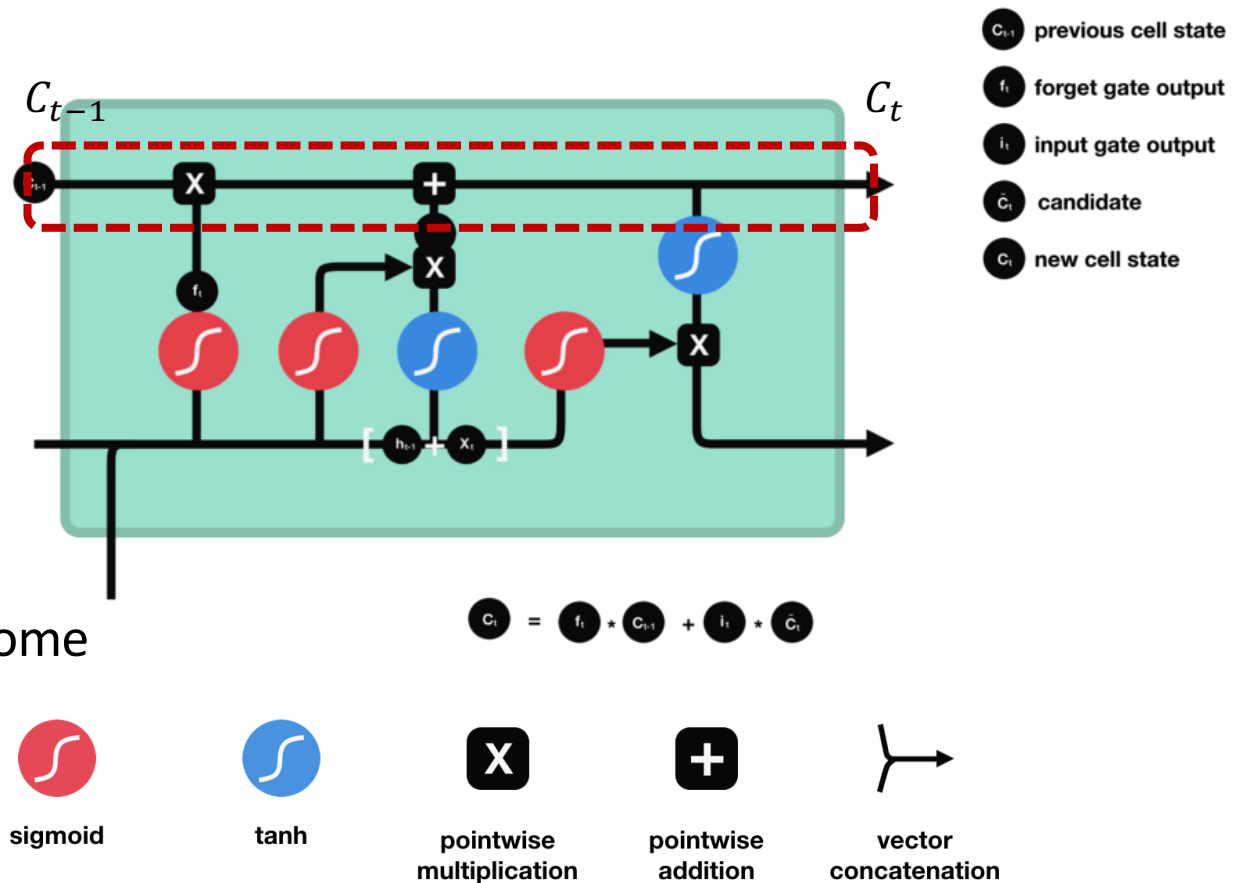
Cell state

- **Cell state:** update the cell state to new values that the network finds relevant.

$$C_t = f_t \circ C_{t-1} + i_t \circ \tilde{C}_t$$

Erase ("forget") some content from the previous state

Keep ("input") some new cell content



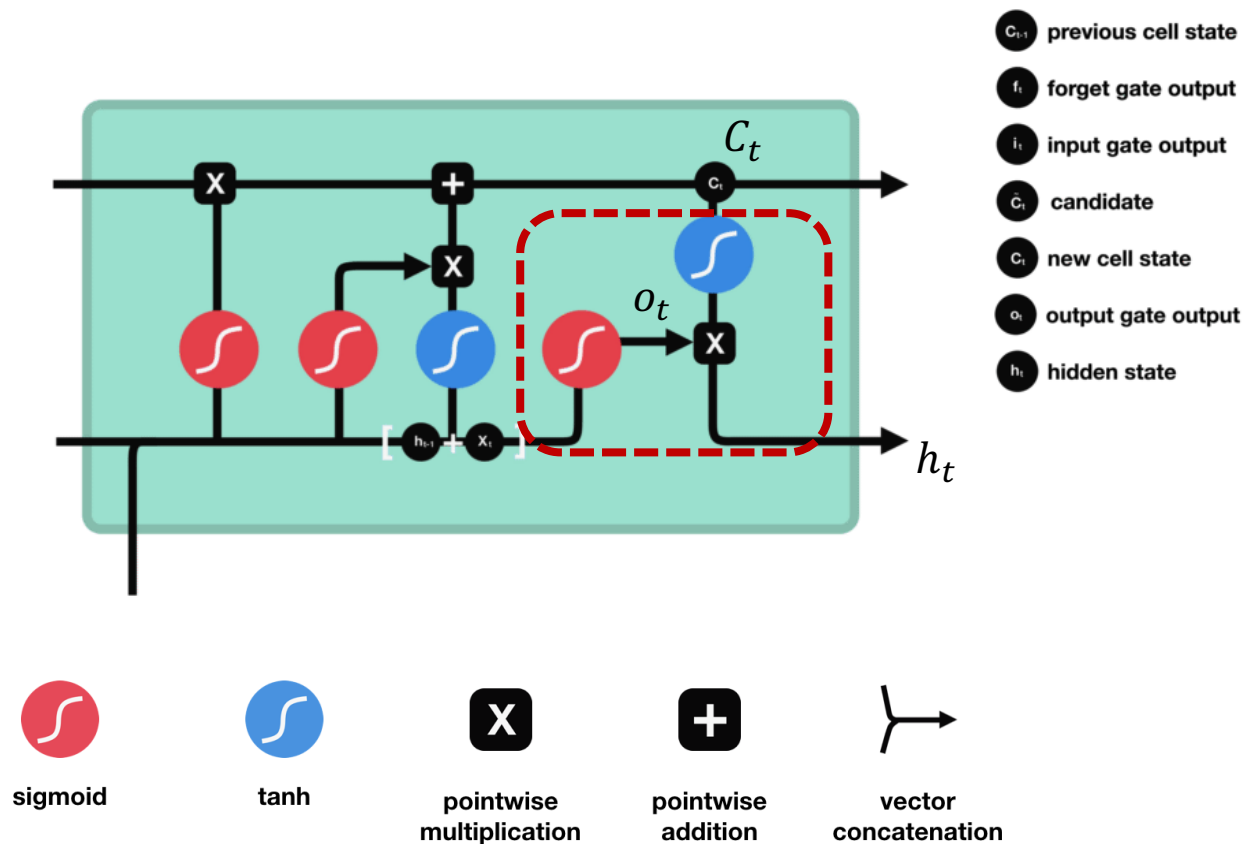
Output gate


- **Output gate:** determines what parts of the cell are output to the hidden state.

$$o_t = \sigma(W_o \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_o)$$

- **Hidden state:** read (“**output**”) some content from the cell.


$$h_t = o_t \circ \tanh(C_t)$$





Vanilla RNN:

$$h_t = \tanh(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_h)$$



LSTM:

$$\begin{aligned} f_t &\in \mathbb{R}^H \\ i_t &\in \mathbb{R}^H \\ o_t &\in \mathbb{R}^H \\ \tilde{C}_t &\in \mathbb{R}^H \\ C_t &\in \mathbb{R}^H \\ h_t &\in \mathbb{R}^H \\ x_t &\in \mathbb{R}^M \\ W_f, W_i, W_o, W_C &\in \mathbb{R}^{H \times (H+M)} \end{aligned}$$

$$f_t = \sigma(W_f \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_f)$$

$$i_t = \sigma(W_i \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_i)$$

$$o_t = \sigma(W_o \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_o)$$

$$\tilde{C}_t = \tanh(W_C \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} + b_C)$$

$$C_t = f_t \circ C_{t-1} + i_t \circ \tilde{C}_t$$

$$h_t = o_t \circ \tanh(C_t)$$

Reference for Topic 1

- Blog by Colah: Understanding LSTM Networks.
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- Blog by Michael Phi: Illustrated Guide to LSTM's and GRU's: A step to step explanation. <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>
- Video lecture by Alexander Amini: MIT course on Recurrent Neural Networks, YouTube.
- Lectures from University of Michigan:
<https://web.eecs.umich.edu/~justincj/teaching/eecs498/FA2020/schedule.html>

Topic 2: LSTM-2

LSTM Architecture

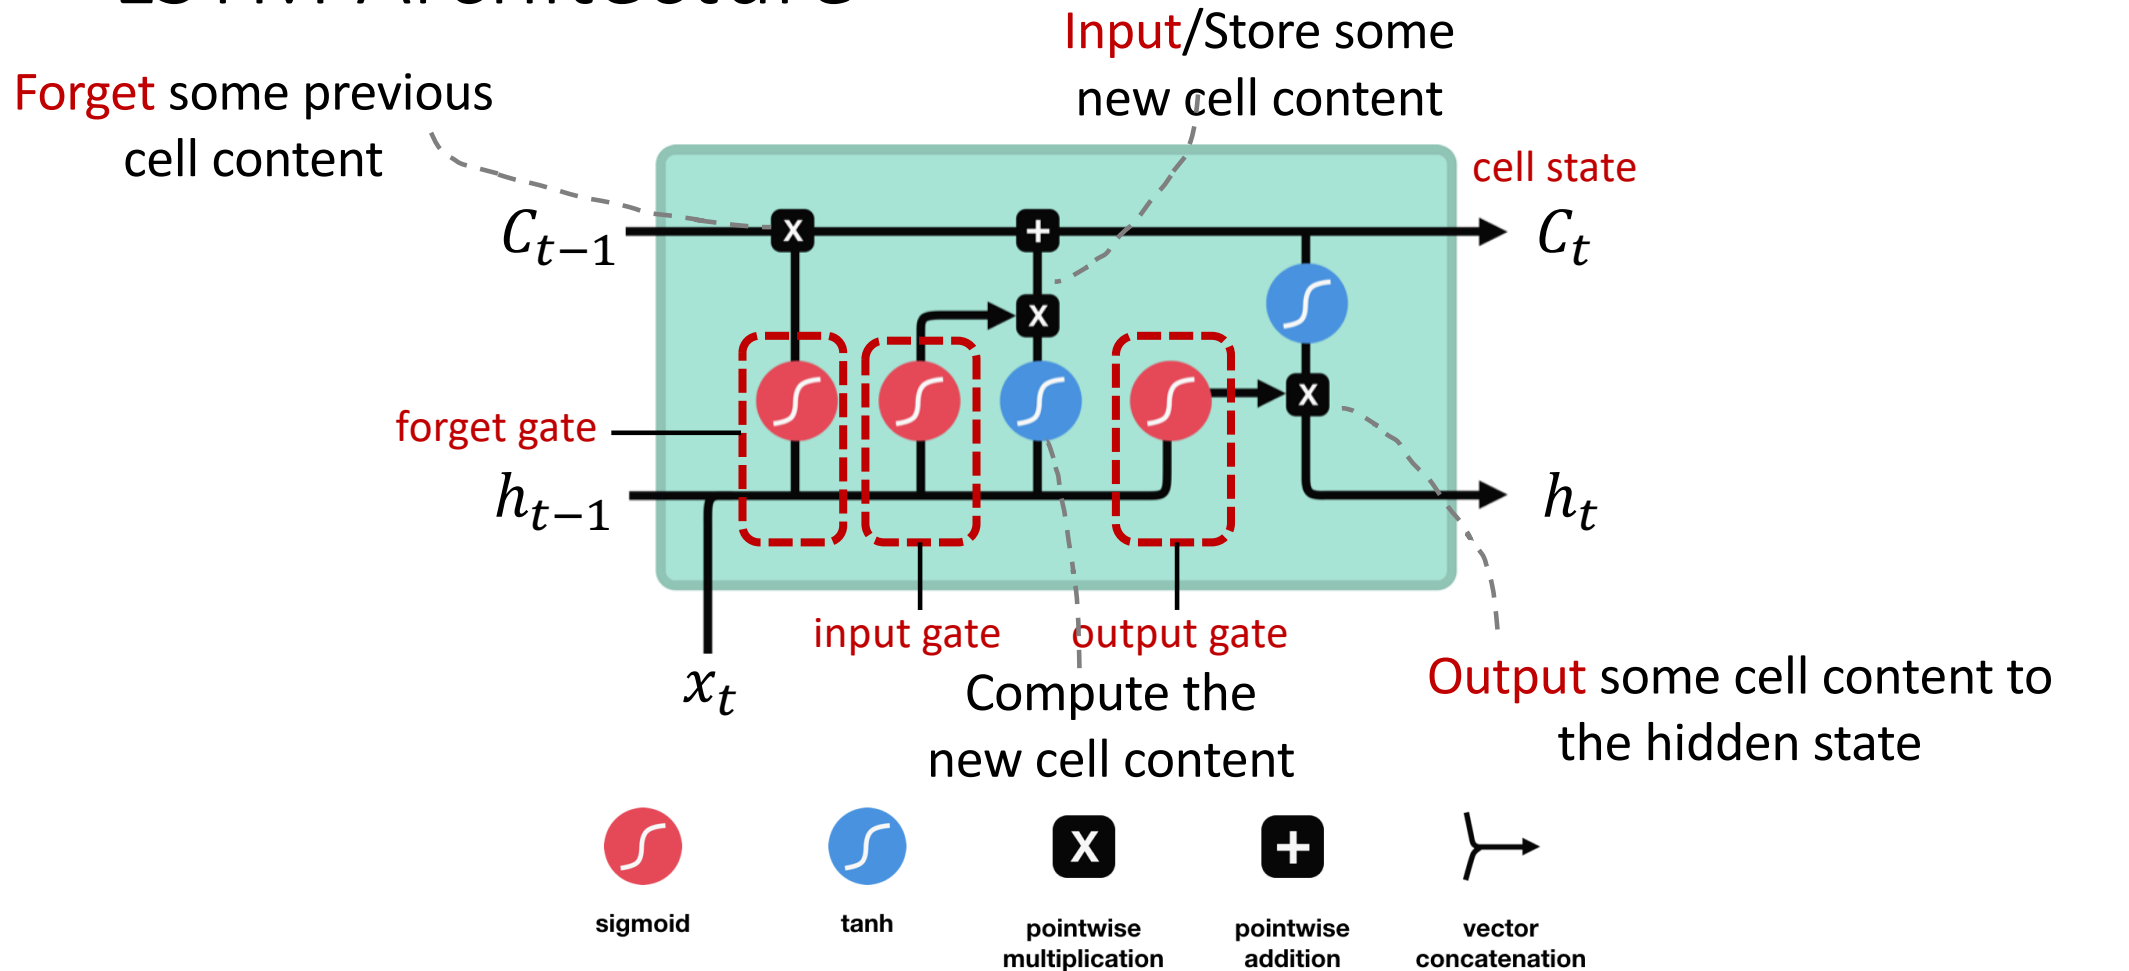



Image from: <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

LSTM Gradient Flow



LSTMs solve the vanishing/exploding gradient problem
using an additive gradient structure.

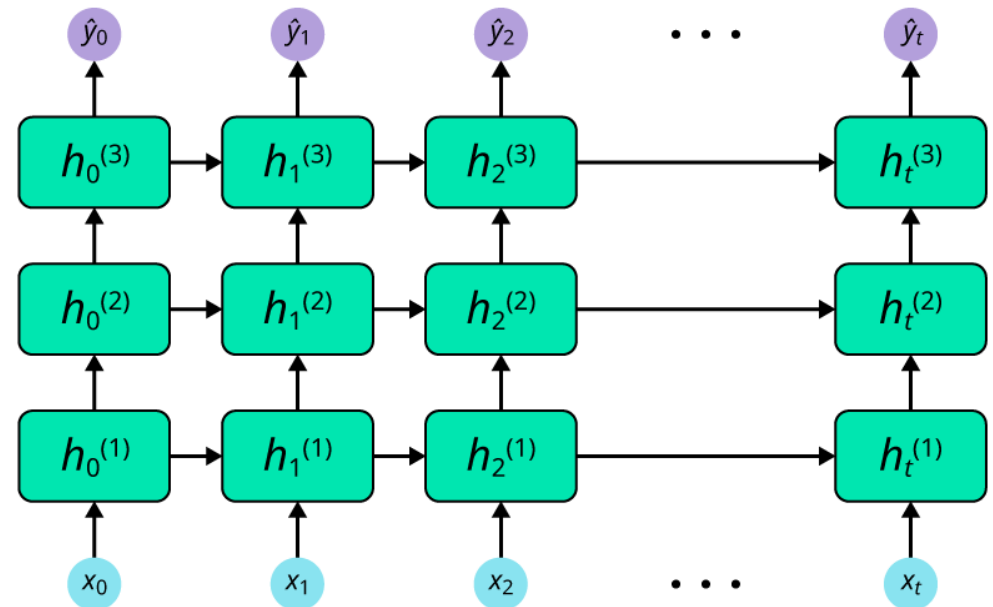
A detailed explanation: <https://medium.com/datadriveninvestor/how-do-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577>

Advanced use of RNNs/LSTMs

- **Multi-layer RNNs (Deep RNNs):**
stack more than one RNN. It increases the representation power of the network, at the cost of higher computational loads.


```
from keras.layers import LSTM
...
model.add(LSTM(32), return_sequences=True)
model.add(LSTM(32), return_sequences=True)
model.add(LSTM(32))
...
```

`True` means: output all the hidden states



Advanced use of RNNs/LSTMs


- **Bidirectional RNNs:** process a sequence in both directions, capturing patterns that may be missed by the chronological-order version alone.



```
from keras.layers import Bidirectional, LSTM
...
model.add(Bidirectional(LSTM(32)))
...
```


Example Tasks: Machine Translation

- Encoder-Decoder [Cho et al, 2014]



Problem: **Encoding bottleneck**
One solution: **Attention Based Encoder-Decoder**

Example Tasks: Image Caption Generation

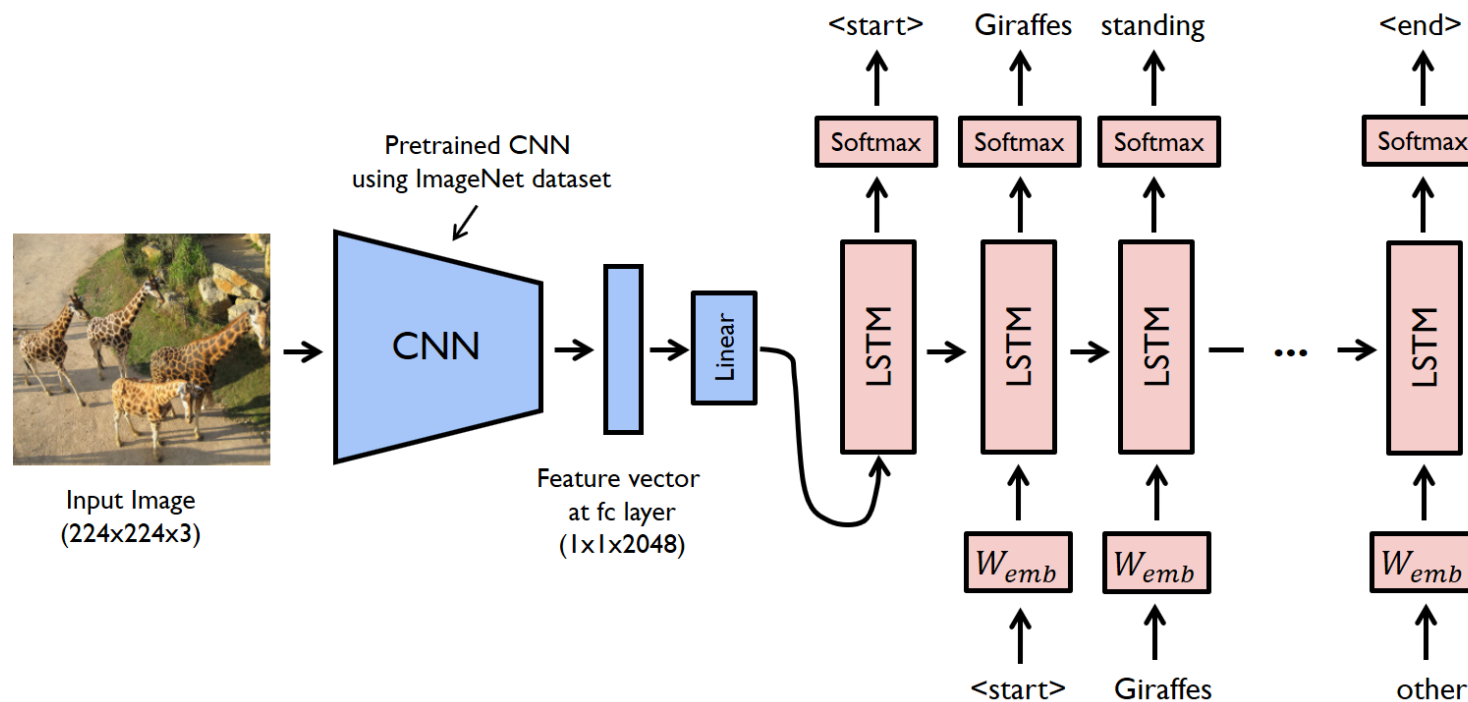


Image from: <https://www.analyticsvidhya.com/blog/2018/04/solving-an-image-captioning-task-using-deep-learning/>

Summary of LSTMs

- Beside the hidden state, also maintain a **cell state** to store **long-term information**.
- Use **gates** to control the flow of information
 - **Forget** gate: gets rid of irrelevant **old** information
 - **Input** gate: stores relevant information from **current** input
 - **Output** gate: **output** a filtered version of the cell state
- Backpropagation through time with **uninterrupted gradient flow**, to avoid the vanishing/exploding gradient problem.

Reference for Topic 2

- Blog by Colah: Understanding LSTM Networks.
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- Blog by Michael Phi: Illustrated Guide to LSTM's and GRU's: A step to step explanation. <https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>
- Video lecture by Alexander Amini: MIT course on Recurrent Neural Networks, YouTube.
- Lectures from University of Michigan:
<https://web.eecs.umich.edu/~justincj/teaching/eecs498/FA2020/schedule.html>
- Blogs by Nir Arbel: How do LSTM networks solve the problem of vanishing gradients: <https://medium.com/datadriveninvestor/how-do-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577>